



STOR 320 Data Transformation II

Lecture 5

Yao Li

Department of Statistics and Operations Research

UNC Chapel Hill

Data Transformation II Info

- Finish Reading Chapter 5 and Practice the Code in R4DS
- Covers
 - The Pipe
 - Statistical Summaries
 - Grouped Summaries
 - Helpful Functions
- Builds Off Last Tutorial

The Pipe

- Useful for Combining Multiple Steps of Operations
- Represented by `%>%`
- Reads as “Then”
- Works Like a Composite Function From Algebra

$$f(x) = 3x + 4$$

$$g(x) = 2x$$

$$h = 1$$

$$f(g(h)) = 3(2(1)) + 4 = 10$$



OUT = h %>%

g() %>%
f()

OUT = 10

The Pipe

- Chaining with the Pipe

```
``{r,eval=F}
f2e.pipedream =
  # Acknowledge the Original Data
  flights %>%
  # Input Original Data and Perform Mutations
  transmute(dep_hr=dep_time%%100+(dep_time%%100)/60,
            sched_dep_hr=sched_dep_time%%100+(sched_dep_time%%100)/60,
            arr_hr=arr_time%%100+(arr_time%%100)/60,
            sched_arr_hr=sched_arr_time%%100+(sched_arr_time%%100)/60) %>%

  mutate(dep_delay_hr=dep_hr-sched_dep_hr,
         arr_delay_hr=arr_hr-sched_arr_hr) %>%

  mutate(percent_dep_delay_hr=percent_rank(dep_delay_hr)) %>%
  # Input Modified Data and Filter the observations
  filter(percent_dep_delay_hr<0.1|percent_dep_delay_hr>0.9) %>%

  # Input Modified Data and Sort according to percent_dep_delay_hr
  arrange(desc(percent_dep_delay_hr))
```

dep_hr <dbl>	sched_dep_hr <dbl>	arr_hr <dbl>	sched_arr_hr <dbl>	dep_delay_hr <dbl>	arr_delay_hr <dbl>
23.35000	8.166667	1.583333	10.33333	15.18333	-8.750000
22.95000	7.983333	1.350000	10.43333	14.96667	-9.083333
22.71667	8.500000	1.000000	11.10000	14.21667	-10.100000
23.40000	10.266667	1.233333	12.45000	13.13333	-11.216667
19.35000	6.250000	21.583333	8.70000	13.10000	12.883333

5 rows | 1-6 of 7 columns

The Pipe

Why use



- Avoid nested functions
- Minimize number of local variables
- Easier to add steps in the sequence

Why not to use

- Debug
- Can't handle multiple inputs
- Can't handle complex code structure

summarize()

- Summarizing All Data
- Using Graphics



Both the histogram and the boxplot are made from summary statistics.

(Statistical Transformations in Ch. 3)

summarize()

- Summarizing All Data



```
``{r}  
dep_delay_hr.summary1 = summarize(f3e.pipedream,  
                                   n=n(),  
                                   mean=mean(dep_delay_hr, na.rm=T),  
                                   var=var(dep_delay_hr, na.rm=T),  
                                   sd=sd(dep_delay_hr, na.rm=T))  
  
dep_delay_hr.summary2 =  
  f3e.pipedream %>%  
  summarize(n=n(),  
            min=min(dep_delay_hr),  
            Q1=quantile(dep_delay_hr, 0.25),  
            Q2=quantile(dep_delay_hr, 0.5),  
            Q3=quantile(dep_delay_hr, 0.75),  
            max=max(dep_delay_hr),  
            IQR=Q3-Q1  
  )  
``
```

n	mean	var	sd
326848	0.192752	0.3506166	0.5921289

n	min	Q1	Q2	Q3	max	IQR
326848	-2.316667	-0.0833333	-0.0333333	0.1666667	4.983333	0.25

summarize() with group_by()

- Summarizing Data by Groups
- Using Graphics



summarize() with group_by()

- Summarizing Data by Groups
 - Using Tables

```
```\r\ngroup.summary1 = f3e.pipedream %>%  
 group_by(origin) %>%
 summarize(n=n())


group.summary2 =
 f3e.pipedream %>%
 group_by(origin) %>%
 summarize(n=n(),
 min=min(dep_delay_hr),
 Q1=quantile(dep_delay_hr,0.25),
 Q2=quantile(dep_delay_hr,0.5),
 Q3=quantile(dep_delay_hr,0.75),
 max=max(dep_delay_hr),
 IQR=Q3-Q1,
 nLow=sum(dep_delay_hr<Q1-1.5*IQR),
 propHigh=mean(dep_delay_hr>Q3+1.5*IQR))
```\r\n
```

origin	n
EWR	117209
JFK	108486
LGA	101153

origin	n	min	Q1	Q2	Q3	max	IQR	nL...	propHigh
<chr>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<dbl>
EWR	117209	-0.4166667	-0.06666667	-0.01666667	0.25	4.983333	0.3166667	0	0.1259204
JFK	108486	-2.3166667	-0.08333333	-0.01666667	0.15	4.983333	0.2333333	2	0.1372988
LGA	101153	-0.5500000	-0.10000000	-0.05000000	0.10	4.983333	0.2000000	7	0.1466491

summarize() with group_by()

- Multiple Groups
 - Using Graphics



summarize() with group_by()

- Multiple Groups
 - Using Tables

```
library(tidyverse)

group.summary3 =
  f3e_ninedream %>%
  group_by(origin, carrier) %>%
  summarize(n=n(),
            min=min(dep_delay_hr),
            Q1=quantile(dep_delay_hr,0.25),
            Q2=quantile(dep_delay_hr,0.5),
            Q3=quantile(dep_delay_hr,0.75),
            max=max(dep_delay_hr),
            IQR=Q3-Q1,
            nLow=sum(dep_delay_hr<Q1-1.5*IQR),
            propHigh=mean(dep_delay_hr>Q3+1.5*IQR))
```

origin	carrier	n	min	Q1	Q2	Q3	max
EWR	9E	1199	-0.2666667	-0.1166667	-0.08333333	-0.01666667	4.416667
EWR	AA	3376	-0.2500000	-0.1000000	-0.05000000	0.05000000	4.750000
EWR	AS	712	-0.3500000	-0.1166667	-0.05000000	0.05000000	3.750000
EWR	B6	6446	-0.3833333	-0.1166667	-0.05000000	0.15000000	4.850000
EWR	DL	4281	-0.2666667	-0.08333333	-0.03333333	0.06666667	4.966667
EWR	EV	41592	-0.4166667	-0.08333333	-0.01666667	0.43333333	4.966667
EWR	MQ	2095	-0.3000000	-0.1000000	-0.03333333	0.30000000	4.950000
EWR	OO	6	-0.1500000	-0.1250000	-0.01666667	0.1791667	2.183333
EWR	UA	45561	-0.3000000	-0.0500000	0.00000000	0.20000000	4.966667
EWR	US	4326	-0.3166667	-0.1000000	-0.06666667	0.00000000	4.166667
EWR	VX	1554	-0.3333333	-0.08333333	-0.02500000	0.08333333	4.916667
EWR	WN	6061	-0.2000000	-0.03333333	0.01666667	0.30000000	4.983333
JFK	9E	13801	-0.4000000	-0.08333333	-0.01666667	0.33333333	4.950000
JFK	AA	13617	-0.2500000	-0.06666667	-0.03333333	0.08333333	4.900000
JFK	B6	41005	-0.7166667	-0.06666667	-0.01666667	0.20000000	4.966667
JFK	DL	20551	-0.3000000	-0.06666667	-0.03333333	0.08333333	4.983333
JFK	EV	1315	-0.3166667	-0.1000000	-0.05000000	0.25000000	4.916667

Useful Summary Functions

- Measures of Center
 - `mean()`
 - `median()`
 - `mode()`
- Measures of Spread
 - `var()`
 - `sd()`
 - `IQR()`
 - `mad()`
- Measures of Rank
 - `min()`
 - `max()`
 - `quantile()`

Useful Summary Functions

- Measures of Position
 - Order Matters
 - `first() = x[1]`
 - `last() = x[length(x)]`
 - `nth(,k) = x[k]`
- Counts
 - `n()`
 - `n_distinct()`
- Counts/Proportions for Logical
 - `sum()`
 - `mean()`
 - Example
 - `sum(x>10)`
 - `mean(x>10)`

Case Study

- Flight Accuracy
 - Accurate Flight Means
 - Departure Delay = 0
 - Arrival Delay = 0

- Bad Metric

$$Accuracy = delay_{dep} + delay_{arr}$$

$$Accuracy = (delay_{dep} + delay_{arr})/2$$

- Good Metrics

$$Accuracy = |delay_{dep}| + |delay_{arr}|$$

$$Accuracy = \sqrt{delay_{dep}^2 + delay_{arr}^2}$$

Case Study

- Summary Table
 - Step 1: Accuracy Variable
 - Step 2: Grouping
 - Step 3: Summarize Info
 - Mean
 - Standard Error
 - Lower Bound (95% CI)
 - Upper Bound (95% CI)

```
```{r}
accuracy<-
 f.pipedream3 %>%
 transmute(carrier,origin,
 accuracy=abs(dep_delay_hr)+abs(arr_delay_hr)) %>%
 group_by(carrier,origin) %>%
 summarize(n=n(),
 avg=mean(accuracy,na.rm=T),
 se=sd(accuracy,na.rm=T)/sqrt(n),
 low=avg-2*se,
 high=avg+2*se
)
```
```

Case Study

- Sorted by Average Accuracy
 - Best Carriers/Origin


```
> head(arrange(accuracy, avg), 5)
# A tibble: 5 x 7
# Groups:   carrier [3]
  carrier origin      n    avg      se    low    high
  <chr>    <chr> <int> <dbl> <dbl> <dbl> <dbl>
1 US      EWR    4322 0.505 0.0123 0.481 0.530
2 US      JFK    2960 0.509 0.0152 0.479 0.539
3 US      LGA   12517 0.544 0.0121 0.520 0.569
4 HA      JFK     342 0.556 0.0362 0.483 0.628
5 UA      JFK    4367 0.591 0.0173 0.556 0.625
```

- Worst Carriers/Origin

```
> head(arrange(accuracy, desc(avg)), 5)
# A tibble: 5 x 7
# Groups:   carrier [4]
  carrier origin      n    avg      se    low    high
  <chr>    <chr> <int> <dbl> <dbl> <dbl> <dbl>
1 OO      EWR        6 1.14  0.737 -0.334 2.61
2 EV      LGA    8086 0.986 0.0265 0.933 1.04
3 YV      LGA    542 0.954 0.0597 0.835 1.07
4 FL      LGA   3136 0.952 0.0545 0.843 1.06
5 EV      EWR   40571 0.952 0.0125 0.927 0.977
```


Case Study

- 95% Confidence Intervals



Carrier “OO” Creates a Visual Problem Due to Small Sample Size

Case Study

```
{r}  
ggplot(filter(accuracy, carrier!="oo")) +  
  geom_point(aes(x=carrier, y=avg)) +  
  geom_errorbar(aes(x=carrier, ymin=low, ymax=high)) +  
  facet_grid(origin~.)  
}
```

- 95% Confidence Intervals

