STOR 320 Data Transformation II

Lecture 5

Yao Li

Department of Statistics and Operations Research

UNC Chapel Hill

Data Transformation II Info

- Finish Reading Chapter 5 and Practice the Code in R4DS

- Covers
  - The Pipe
  - Statistical Summaries
  - Grouped Summaries
  - Helpful Functions

- Builds Off Last Tutorial

The Pipe

- Useful for Combining Multiple Steps of Operations

- Represented by %>%

- Reads as "Then"

- Works Like a Composite Function From Algebra

$f(x) = 3x + 4$
$g(x) = 2x$
$h = 1$

$f\big(g(h)\big) = 3\big(2(1)\big) + 4 = 10$

OUT = h %>%
       g() %>%
          f()

OUT = 10

The Pipe

```{r,eval=F}
f2e.pipedream =
  # Acknowledge the Original Data
  flights %>%
  # Input Original Data and Perform Mutations
  transmute(dep_hr=dep_time%/%100+(dep_time%%100)/60,
            sched_dep_hr=sched_dep_time%/%100+(sched_dep_time%%100)/60,
            arr_hr=arr_time%/%100+(arr_time%%100)/60,
            sched_arr_hr=sched_arr_time%/%100+(sched_arr_time%%100)/60) %>%

  mutate(dep_delay_hr=dep_hr-sched_dep_hr,
         arr_delay_hr=arr_hr-sched_arr_hr) %>%

  mutate(percent_dep_delay_hr=percent_rank(dep_delay_hr)) %>%
  # Input Modified Data and Filter the observations
  filter(percent_dep_delay_hr<0.1|percent_dep_delay_hr>0.9) %>%

  # Input Modified Data and Sort according to percent_dep_delay_hr
  arrange(desc(percent_dep_delay_hr))
```

| dep_hr <dbl>	sched_dep_hr <dbl>	arr_hr <dbl>	sched_arr_hr <dbl>	dep_delay_hr <dbl>	arr_delay_hr <dbl>
23.35000	8.166667	1.583333	10.33333	15.18333	-8.750000
22.95000	7.983333	1.350000	10.43333	14.96667	-9.083333
22.71667	8.500000	1.000000	11.10000	14.21667	-10.100000
23.40000	10.266667	1.233333	12.45000	13.13333	-11.216667
19.35000	6.250000	21.583333	8.70000	13.10000	12.883333

5 rows | 1–6 of 7 columns

The Pipe

Why use

- Avoid nested functions

- Minimize number of local variables

- Easier to add steps in the sequence

Why not to use

- Debug

- Can't handle multiple inputs

- Can't handle complex code structure

summarize()

- Summarizing All Data

    - Using Graphics

Both the histogram and the boxplot are made from summary statistics.

(**Statistical Transformations** in Ch. 3)

summarize()

- Summarizing All Data

```{r}
dep_delay_hr.summary1 = summarize(f3e.pipedream,
                                  n=n(),
                                  mean=mean(dep_delay_hr,na.rm=T),
                                  var=var(dep_delay_hr,na.rm=T),
                                  sd=sd(dep_delay_hr,na.rm=T))

dep_delay_hr.summary2 =
  f3e.pipedream %>%
  summarize(n=n(),
            min=min(dep_delay_hr),
            Q1=quantile(dep_delay_hr,0.25),
            Q2=quantile(dep_delay_hr,0.5),
            Q3=quantile(dep_delay_hr,0.75),
            max=max(dep_delay_hr),
            IQR=Q3-Q1
  )
```

n	mean	var	sd
326848	0.192752	0.3506166	0.5921289

n	min	Q1	Q2	Q3	max	IQR
326848	-2.316667	-0.0833333	-0.0333333	0.1666667	4.983333	0.25

7

summarize() with group_by()

- Summarizing Data by Groups

  - Using Graphics

summarize() with group_by()

- Summarizing Data by Groups

  - Using Tables

```{r}
group.summary1 = f3e.pipedream %>%
                 group_by(origin) %>%
                 summarize(n=n())

group.summary2 =
  f3e.pipedream %>%
  group_by(origin) %>%
  summarize(n=n(),
            min=min(dep_delay_hr),
            Q1=quantile(dep_delay_hr,0.25),
            Q2=quantile(dep_delay_hr,0.5),
            Q3=quantile(dep_delay_hr,0.75),
            max=max(dep_delay_hr),
            IQR=Q3-Q1,
            nLow=sum(dep_delay_hr<Q1-1.5*IQR),
            propHigh=mean(dep_delay_hr>Q3+1.5*IQR)
  )
```
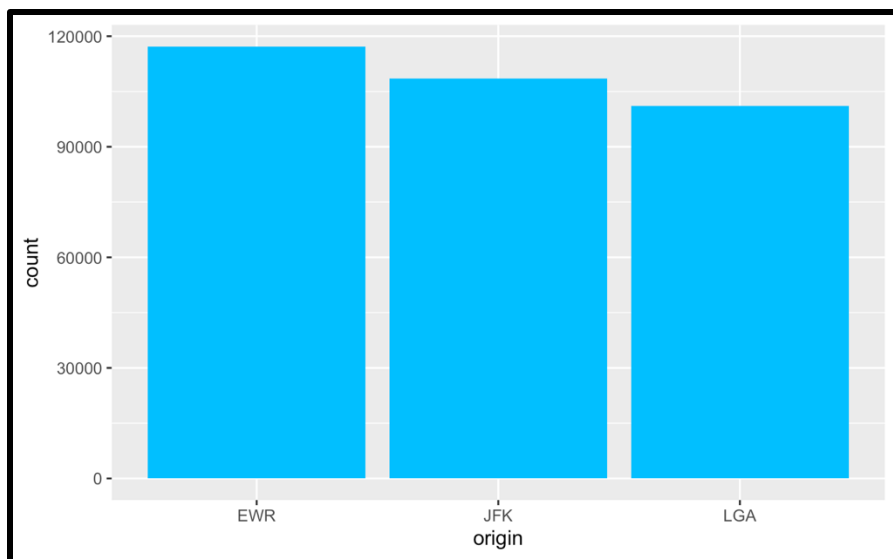
origin	n
EWR	117209
JFK	108486
LGA	101153

origin <chr>	n <int>	min <dbl>	Q1 <dbl>	Q2 <dbl>	Q3 <dbl>	max <dbl>	IQR <dbl>	nL... <int>	propHigh <dbl>
EWR	117209	−0.4166667	−0.06666667	−0.01666667	0.25	4.983333	0.3166667	0	0.1259204
JFK	108486	−2.3166667	−0.08333333	−0.01666667	0.15	4.983333	0.2333333	2	0.1372988
LGA	101153	−0.5500000	−0.10000000	−0.05000000	0.10	4.983333	0.2000000	7	0.1466491

summarize() with group_by()

- Multiple Groups

  - Using Graphics

summarize() with group_by()

- Multiple Groups

  - Using Tables

```{r}
group.summary3 =
  f3e_pipedream %>%
  group_by(origin,carrier) %>%
  summarize(n=n(),
            min=min(dep_delay_hr),
            Q1=quantile(dep_delay_hr,0.25),
            Q2=quantile(dep_delay_hr,0.5),
            Q3=quantile(dep_delay_hr,0.75),
            max=max(dep_delay_hr),
            IQR=Q3-Q1,
            nLow=sum(dep_delay_hr<Q1-1.5*IQR),
            propHigh=mean(dep_delay_hr>Q3+1.5*IQR)
  )
```

origin	carrier	n	min	Q1	Q2	Q3	max
EWR	9E	1199	-0.2666667	-0.1166667	-0.0833333	-0.0166667	4.416667
EWR	AA	3376	-0.2500000	-0.1000000	-0.0500000	0.0500000	4.750000
EWR	AS	712	-0.3500000	-0.1166667	-0.0500000	0.0500000	3.750000
EWR	B6	6446	-0.3833333	-0.1166667	-0.0500000	0.1500000	4.850000
EWR	DL	4281	-0.2666667	-0.0833333	-0.0333333	0.0666667	4.966667
EWR	EV	41592	-0.4166667	-0.0833333	-0.0166667	0.4333333	4.966667
EWR	MQ	2095	-0.3000000	-0.1000000	-0.0333333	0.3000000	4.950000
EWR	OO	6	-0.1500000	-0.1250000	-0.0166667	0.1791667	2.183333
EWR	UA	45561	-0.3000000	-0.0500000	0.0000000	0.2000000	4.966667
EWR	US	4326	-0.3166667	-0.1000000	-0.0666667	0.0000000	4.166667
EWR	VX	1554	-0.3333333	-0.0833333	-0.0250000	0.0833333	4.916667
EWR	WN	6061	-0.2000000	-0.0333333	0.0166667	0.3000000	4.983333
JFK	9E	13801	-0.4000000	-0.0833333	-0.0166667	0.3333333	4.950000
JFK	AA	13617	-0.2500000	-0.0666667	-0.0333333	0.0833333	4.900000
JFK	B6	41005	-0.7166667	-0.0666667	-0.0166667	0.2000000	4.966667
JFK	DL	20551	-0.3000000	-0.0666667	-0.0333333	0.0833333	4.983333
JFK	EV	1315	-0.3166667	-0.1000000	-0.0500000	0.2500000	4.916667

Useful Summary Functions

- Measures of Center
  - mean()
  - median()
  - mode()

- Measures of Spread
  - var()
  - sd()
  - IQR()
  - mad()

- Measures of Rank
  - min()
  - max()
  - quantile()

Useful Summary Functions

- Measures of Position
  - Order Matters
  - first() = x[1]
  - last() = x[length(x)]
  - nth(,k) = x[k]

- Counts
  - n()
  - n_distinct()

- Counts/Proportions for Logical
  - sum()
  - mean()
  - Example
    - sum(x>10)
    - mean(x>10)

Case Study

- Flight Accuracy

    - Accurate Flight Means
        - Departure Delay = 0
        - Arrival Delay = 0

    - Bad Metric

    $$Accuracy = delay_{dep} + delay_{arr}$$
    $$Accuracy = (delay_{dep} + delay_{arr})/2$$

- Good Metrics

    $$Accuracy = |delay_{dep}| + |delay_{arr}|$$
    $$Accuracy = \sqrt{delay_{dep}^2 + delay_{arr}^2}$$

Case Study

- Summary Table
  - Step 1: Accuracy Variable
  - Step 2: Grouping
  - Step 3: Summarize Info
    - Mean
    - Standard Error
    - Lower Bound (95% CI)
    - Upper Bound (95% CI)

```{r}
accuracy<-
  f.pipedream3 %>%
  transmute(carrier,origin,
    accuracy=abs(dep_delay_hr)+abs(arr_delay_hr)) %>%
  group_by(carrier,origin) %>%
  summarize(n=n(),
    avg=mean(accuracy,na.rm=T),
    se=sd(accuracy,na.rm=T)/sqrt(n),
    low=avg-2*se,
    high=avg+2*se
  )
```

Case Study

- Sorted by Average Accuracy
  - Best Carriers/Origin

```
> head(arrange(accuracy,avg),5)
# A tibble: 5 x 7
# Groups:   carrier [3]
  carrier origin       n   avg       se    low  high
  <chr>   <chr>    <int> <dbl>    <dbl>  <dbl> <dbl>
1 US      EWR       4322 0.505 0.0123  0.481 0.530
2 US      JFK       2960 0.509 0.0152  0.479 0.539
3 US      LGA      12517 0.544 0.0121  0.520 0.569
4 HA      JFK        342 0.556 0.0362  0.483 0.628
5 UA      JFK       4367 0.591 0.0173  0.556 0.625
```

- Worst Carriers/Origin

```
> head(arrange(accuracy,desc(avg)),5)
# A tibble: 5 x 7
# Groups:   carrier [4]
  carrier origin       n   avg       se    low  high
  <chr>   <chr>    <int> <dbl>    <dbl>  <dbl> <dbl>
1 OO      EWR          6 1.14  0.737  -0.334 2.61
2 EV      LGA       8086 0.986 0.0265  0.933 1.04
3 YV      LGA        542 0.954 0.0597  0.835 1.07
4 FL      LGA       3136 0.952 0.0545  0.843 1.06
5 EV      EWR      40571 0.952 0.0125  0.927 0.977
```

16

Case Study

- 95% Confidence Intervals

Carrier "OO" Creates a Visual Problem Due to Small Sample Size

Case Study

```r
ggplot(filter(accuracy,carrier!="oo")) +
  geom_point(aes(x=carrier,y=avg)) +
  geom_errorbar(aes(x=carrier,ymin=low,ymax=high)) +
  facet_grid(origin~.)
```

- 95% Confidence Intervals