



STOR 320 Data Transformation I

Lecture 3

Yao Li

Department of Statistics and Operations Research

UNC Chapel Hill

Introduction

- Read Chapter 5
- Goal: Their Data → Your Data
- Covers:
 - Data Subsetting
 - Data Ordering
 - Variable Selecting
 - Variable Creating
- Help: dplyr Package in R

NYC Flights Meta Data

- Requirements:

```
> install.packages(nycflights13)  
> library(nycflights13)
```

- All 2013 Flights from NYC
- US Bureau of Trans. Statistics
- To View all Data, Use `> View(flights)`
- For more information, `> ?flights`

NYC Flights Data

4

```
> flights
# A tibble: 336,776 x 19
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1     1     517           515           2     830           819
2  2013     1     1     533           529           4     850           830
3  2013     1     1     542           540           2     923           850
4  2013     1     1     544           545          -1    1004          1022
5  2013     1     1     554           600          -6     812           837
6  2013     1     1     554           558          -4     740           728
7  2013     1     1     555           600          -5     913           854
8  2013     1     1     557           600          -3     709           723
9  2013     1     1     557           600          -3     838           846
10 2013     1     1     558           600          -2     753           745
# ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

- Four Different Types of Variables
 - int = integer
 - dbl = double
 - chr = character
 - dtm = date and times
- Other Types of Variables
 - lgl = logical (TRUE or FALSE)
 - fctr = factor
 - date = dates

Basics of dplyr: 5 Key Functions

- 5 Key Functions
 - `filter()` = Chooses Observations Based on Values
 - `arrange()` = Sorts Observations
 - `select()` = Chooses Variables
 - `mutate()` = Creates New Variables
 - `summarise()` = Generates Statistics From Data

Basics of dplyr

- Function Usage
 - First, Specify the Dataset
 - Next, Specify What to Do with the Data
 - Result is a New Dataset

```
> filter(flights, month==9)
```

year <int>	month <int>	day <int>	dep_time <int>	sched_dep_time <int>	dep_delay <dbl>	arr_time <int>	sched_arr_time <int>
2013	9	1	9	2359	10	343	340
2013	9	1	117	2245	152	218	2359
2013	9	1	508	516	-8	717	800
2013	9	1	537	545	-8	849	855
2013	9	1	537	545	-8	906	921
2013	9	1	549	600	-11	815	850
2013	9	1	552	600	-8	843	905
2013	9	1	553	600	-7	809	834
2013	9	1	554	600	-6	700	716
2013	9	1	554	600	-6	803	823

1-10 of 27,574 rows | 1-8 of 19 columns

Previous 1 2 3 4 5 6 ... 100 Next

Comparisons

- Important Operators
 - Less Than (<)
 - Greater Than (>)
 - Not Equal (!=)
 - Equal (==)
- Returns TRUE or FALSE

```
{r}
x=3
y=4
x<y
x>y
x!=y
x==y
```

```
[1] TRUE
[1] FALSE
[1] TRUE
[1] FALSE
```

Numerical Precision

- Problem

```
> x=1/49  
> y=49  
> x*y==1  
[1] FALSE  
> near(x*y,1)  
[1] TRUE
```

- Solution

```
> x*y  
[1] 1  
> near(x*y,1)  
[1] TRUE
```


Logical Operators

- Boolean Logic
 - And (&)
 - Or (|)
 - Not (!)

- Example 

```
> filter(flights, month==9&day==1)
```

```
> x = TRUE
> y = FALSE
> # Basic
> x&y
[1] FALSE
> x|y
[1] TRUE
> !x
[1] FALSE
>
> # Combined
> !x|!y
[1] TRUE
> !(x&y)
[1] TRUE
> !x&!y
[1] FALSE
```

Missing Values

- Represented by NA
 - Enduring Questions
 - To Impute or Not Impute
 - To Ignore or Not Ignore
 - Handling Should Be Explained
 - Be Careful When Performing Operations on Missing Data



```
NA > 5
#> [1] NA

10 == NA
#> [1] NA

NA + 10
#> [1] NA

NA / 2
#> [1] NA
```

Missing Values

```
> male.age=c(NA,20,21,35,22,NA)
> female.age=c(21,NA,23,33,22,NA)
> age.data=tibble(ma=male.age,fa=female.age)
> age.data
# A tibble: 6 x 2
      ma     fa
  <dbl> <dbl>
1    NA    21
2    20    NA
3    21    23
4    35    33
5    22    22
6    NA    NA
>
> is.na(male.age)
[1] TRUE FALSE FALSE FALSE FALSE TRUE
> na.omit(age.data)
# A tibble: 3 x 2
      ma     fa
  <dbl> <dbl>
1    21    23
2    35    33
3    22    22
> mean(male.age)
[1] NA
> mean(male.age,na.rm=T)
[1] 24.5
```

filter()

- Used to Subset Observations Based on Their Values
 - Selects Row if TRUE
 - Removes Row if FALSE
- Examples:
 - All Flights from 9/13/2018 Out of LaGuardia Airport

```
> filter(flights, month==9,day==13,origin =="LGA")
```

- All Dec. or Nov. Flights

```
> filter(flights, month==11|month==12)
```

```
> filter(flights, month %in% c(11,12))
```

filter()

- Examples:
 - Don't Want Flights with Unusual Delays (> 120 min.)

```
> filter(flights, !(arr_delay>120 | dep_delay>120) )
```

```
> filter(flights, arr_delay <= 120, dep_delay <= 120)
```

- Want Flights with No Delays

```
> filter(flights, dep_delay==0, arr_delay==0)
```

```
> filter(flights, dep_delay==0 & arr_delay==0)
```

$!(x \& y) = !x \mid !y$

filter()

- Examples:

- Want Flights Missing Air Time

```
> filter(flights, is.na(air_time) )
```

- Do not Want Flights Missing Air Time

```
> filter(flights, !is.na(air_time) )
```

- Remove All Cases with Missing Values

```
> na.omit(flights)
```

arrange()

- Used to Sort Observations
- Sort flights by date

```
```{r}  
arrange(flights, year, month, day)
```
```

| year
<int> | month
<int> | day
<int> | dep_time
<int> | sched_dep_time
<int> | dep_delay
<dbl> | arr_time
<int> |
|---------------|----------------|--------------|-------------------|-------------------------|--------------------|-------------------|
| 2013 | 1 | 1 | 517 | 515 | 2 | 830 |
| 2013 | 1 | 1 | 533 | 529 | 4 | 850 |
| 2013 | 1 | 1 | 542 | 540 | 2 | 923 |
| 2013 | 1 | 1 | 544 | 545 | -1 | 1004 |
| 2013 | 1 | 1 | 554 | 600 | -6 | 812 |
| 2013 | 1 | 1 | 554 | 558 | -4 | 740 |
| 2013 | 1 | 1 | 555 | 600 | -5 | 913 |
| 2013 | 1 | 1 | 557 | 600 | -3 | 709 |
| 2013 | 1 | 1 | 557 | 600 | -3 | 838 |
| 2013 | 1 | 1 | 558 | 600 | -2 | 753 |

arrange()

- Sorting Experiment

```
{r}  
head(arrange(flights, day, dep_time))
```

| year | month | day | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time |
|-------|-------|-------|----------|----------------|-----------|----------|----------------|
| <int> | <int> | <int> | <int> | <int> | <dbl> | <int> | <int> |
| 2013 | 7 | 1 | 1 | 2029 | 212 | 236 | 2359 |
| 2013 | 6 | 1 | 2 | 2359 | 3 | 341 | 350 |
| 2013 | 7 | 1 | 2 | 2359 | 3 | 344 | 344 |
| 2013 | 3 | 1 | 4 | 2159 | 125 | 318 | 56 |
| 2013 | 11 | 1 | 5 | 2359 | 6 | 352 | 345 |
| 2013 | 5 | 1 | 9 | 1655 | 434 | 308 | 2020 |

6 rows | 1-8 of 19 columns

```
{r}  
head(arrange(flights, desc(day), dep_time))
```

| year | month | day | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time |
|-------|-------|-------|----------|----------------|-----------|----------|----------------|
| <int> | <int> | <int> | <int> | <int> | <dbl> | <int> | <int> |
| 2013 | 1 | 31 | 1 | 2100 | 181 | 124 | 2225 |
| 2013 | 1 | 31 | 4 | 2359 | 5 | 455 | 444 |
| 2013 | 1 | 31 | 7 | 2359 | 8 | 453 | 437 |
| 2013 | 7 | 31 | 10 | 2359 | 11 | 344 | 340 |
| 2013 | 1 | 31 | 12 | 2250 | 82 | 132 | 7 |
| 2013 | 12 | 31 | 13 | 2359 | 14 | 439 | 437 |

6 rows | 1-8 of 19 columns

arrange()

- Handling NA

```
{r}
miss.data=tibble(x=c(1,1,NA,3,3,NA),
                  y=c(NA,4,NA,5,NA,7))
miss.data|
```

| x | y |
|-------|-------|
| <dbl> | <dbl> |
| 1 | NA |
| 1 | 4 |
| NA | NA |
| 3 | 5 |
| 3 | NA |
| NA | 7 |

6 rows

```
{r}
arrange(miss.data,x|)
```

| x | y |
|-------|-------|
| <dbl> | <dbl> |
| 1 | NA |
| 1 | 4 |
| 3 | 5 |
| 3 | NA |
| NA | NA |
| NA | 7 |

```
{r}
arrange(miss.data,desc(x))|
```

| x | y |
|-------|-------|
| <dbl> | <dbl> |
| 3 | 5 |
| 3 | NA |
| 1 | NA |
| 1 | 4 |
| NA | NA |
| NA | 7 |

select()

- Used to Select Variables
- Why? Not All Variables are Created Equal
- Need to Know Variable Names

```
> names(flights)
[1] "year"          "month"          "day"
[4] "dep_time"      "sched_dep_time" "dep_delay"
[7] "arr_time"      "sched_arr_time" "arr_delay"
[10] "carrier"       "flight"         "tailnum"
[13] "origin"        "dest"           "air_time"
[16] "distance"      "hour"           "minute"
[19] "time_hour"
```

- Basic Examples

select()

- Select Only Year, Month, Day

```
> data1=select(flights,year,month,day)
> names(data1)
[1] "year" "month" "day"
```

- Select All Variables Between dep_time to arr_delay

```
> data2=select(flights,dep_time:arr_delay)
> names(data2)
[1] "dep_time"          "sched_dep_time"
[3] "dep_delay"         "arr_time"
[5] "sched_arr_time"    "arr_delay"
```

- Deselect All Variables Between dep_time to arr_delay

```
> data3=select(flights,-(dep_time:arr_delay))
> names(data3)
[1] "year"      "month"     "day"
[4] "carrier"   "flight"    "tailnum"
[7] "origin"    "dest"      "air_time"
[10] "distance"  "hour"      "minute"
[13] "time_hour"
```

select()

- Select Based on Column Index

```
> length(names(flights))  
[1] 19  
> data4=select(flights,c(1,3,8,12))  
> names(data4)  
[1] "year"  
[2] "day"  
[3] "sched_arr_time"  
[4] "tailnum"
```

- Deselect Based on Column Index

```
> length(names(flights))  
[1] 19  
> data5=select(flights,-c(1,3,8,12))  
> names(data5)  
[1] "month"  
[2] "dep_time"  
[3] "sched_dep_time"  
[4] "dep_delay"  
[5] "arr_time"  
[6] "arr_delay"  
[7] "carrier"  
[8] "flight"  
[9] "origin"  
[10] "dest"  
[11] "air_time"  
[12] "distance"  
[13] "hour"  
[14] "minute"  
[15] "time_hour"
```

select()

- Select Based on Text
 - starts_with("TEXT")

```
> data6=select(flights,starts_with("dep"))  
> names(data6)  
[1] "dep_time" "dep_delay"
```

- ends_with("TEXT")

```
> data7=select(flights,ends_with("delay"))  
> names(data7)  
[1] "dep_delay" "arr_delay"
```

- contains("TEXT")

```
> data8=select(flights,contains("ar"))  
> names(data8)  
[1] "year" "arr_time"  
[3] "sched_arr_time" "arr_delay"  
[5] "carrier"
```

select()

- Renaming Variables
 - Can Use select()

```
> data9=select(flights,yr=year)
> names(data9)
[1] "yr"
```

- But Use rename()

```
> data10=rename(flights,yr=year)
> names(data10)
[1] "yr"           "month"
[3] "day"          "dep_time"
[5] "sched_dep_time" "dep_delay"
[7] "arr_time"     "sched_arr_time"
[9] "arr_delay"    "carrier"
[11] "flight"       "tailnum"
[13] "origin"       "dest"
[15] "air_time"     "distance"
[17] "hour"         "minute"
[19] "time_hour"
```


select()

- Reordering Variables

```
> head(flights)
# A tibble: 6 x 19
  year month   day dep_time sched_dep_time
  <int> <int> <int>   <int>         <int>
1  2013     1     1     517             515
2  2013     1     1     533             529
3  2013     1     1     542             540
4  2013     1     1     544             545
5  2013     1     1     554             600
6  2013     1     1     554             558
# ... with 14 more variables: dep_delay <dbl>,
#   arr_time <int>, sched_arr_time <int>,
#   arr_delay <dbl>, carrier <chr>,
#   flight <int>, tailnum <chr>, origin <chr>,
#   dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
> data11=select(flights,dep_time,arr_time,
+               air_time,everything())
> head(data11)
# A tibble: 6 x 19
  dep_time arr_time air_time year month   day
  <int>    <int>    <dbl> <int> <int> <int>
1     517     830     227  2013     1     1
2     533     850     227  2013     1     1
3     542     923     160  2013     1     1
4     544    1004     183  2013     1     1
5     554     812     116  2013     1     1
6     554     740     150  2013     1     1
# ... with 13 more variables:
#   sched_dep_time <int>, dep_delay <dbl>,
#   sched_arr_time <int>, arr_delay <dbl>,
#   carrier <chr>, flight <int>, tailnum <chr>,
#   origin <chr>, dest <chr>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

mutate()

- Used to Create New Variables
 - Creative New Metrics
 - Modify Units
 - Transform Variables
 - Unique Identifiers
 - Numeric to Categorical
 - Categorical to Numeric
- Reduced Dataset

```
library(tidyverse)
flights_sml <- select(flights, year:day,
                     starts_with("dep"),
                     starts_with("arr"),
                     distance, air_time)
head(flights_sml)
```

| year | month | day | dep_time | dep_delay | arr_time | arr_delay | distance | air_time |
|-------|-------|-------|----------|-----------|----------|-----------|----------|----------|
| <int> | <int> | <int> | <int> | <dbl> | <int> | <dbl> | <dbl> | <dbl> |
| 2013 | 1 | 1 | 517 | 2 | 830 | 11 | 1400 | 227 |
| 2013 | 1 | 1 | 533 | 4 | 850 | 20 | 1416 | 227 |
| 2013 | 1 | 1 | 542 | 2 | 923 | 33 | 1089 | 160 |
| 2013 | 1 | 1 | 544 | -1 | 1004 | -18 | 1576 | 183 |
| 2013 | 1 | 1 | 554 | -6 | 812 | -25 | 762 | 116 |
| 2013 | 1 | 1 | 554 | -4 | 740 | 12 | 719 | 150 |

mutate()

- Example of mutate()

```
{r}
mutate_flights_sml<-mutate(flights_sml,
                           gain=arr_delay-dep_delay,
                           speed=distance/air_time*60)
head(select(mutate_flights_sml,gain,speed,everything()))
```

| gain
<dbl> | speed
<dbl> | year
<int> | month
<int> | day
<int> | dep_time
<int> | dep_delay
<dbl> | arr_time
<int> | arr_delay
<dbl> |
|---------------|----------------|---------------|----------------|--------------|-------------------|--------------------|-------------------|--------------------|
| 9 | 370.0441 | 2013 | 1 | 1 | 517 | 2 | 830 | 11 |
| 16 | 374.2731 | 2013 | 1 | 1 | 533 | 4 | 850 | 20 |
| 31 | 408.3750 | 2013 | 1 | 1 | 542 | 2 | 923 | 33 |
| -17 | 516.7213 | 2013 | 1 | 1 | 544 | -1 | 1004 | -18 |
| -19 | 394.1379 | 2013 | 1 | 1 | 554 | -6 | 812 | -25 |
| 16 | 287.6000 | 2013 | 1 | 1 | 554 | -4 | 740 | 12 |

- Example of transmute()

```
{r}
transmute_flights_sml<-transmute(flights_sml,
                                  gain=arr_delay-dep_delay,
                                  speed=distance/air_time*60)
head(select(transmute_flights_sml,gain,speed,everything()))
```

| gain
<dbl> | speed
<dbl> |
|---------------|----------------|
| 9 | 370.0441 |
| 16 | 374.2731 |
| 31 | 408.3750 |
| -17 | 516.7213 |
| -19 | 394.1379 |
| 16 | 287.6000 |

mutate()

- Plethora of Examples
 - Basic and Modular Arithmetic

dep_time, arr_time

Actual departure and arrival times (format HHMM or HMM), local tz.

```
{r}
flights1=transmute(flights,
  dep_time,
  hour=dep_time%%100,
  minute=dep_time%%100)
flights1
```

| dep_time
<int> | hour
<dbl> | minute
<dbl> |
|-------------------|---------------|-----------------|
| 517 | 5 | 17 |
| 533 | 5 | 33 |
| 542 | 5 | 42 |

$$\begin{aligned} 517 &= 100 * 5 + 17 \\ &= 100 * (517 \% 100) + (517 \% 100) \end{aligned}$$

```
{r}
flights2=transmute(flights1,
  dep_time,
  hour,
  minute,
  hrs_since_midnight=hour+minute/60)
flights2
```

| dep_time
<int> | hour
<dbl> | minute
<dbl> | hrs_since_midnight
<dbl> |
|-------------------|---------------|-----------------|-----------------------------|
| 517 | 5 | 17 | 5.283333 |
| 533 | 5 | 33 | 5.550000 |
| 542 | 5 | 42 | 5.700000 |


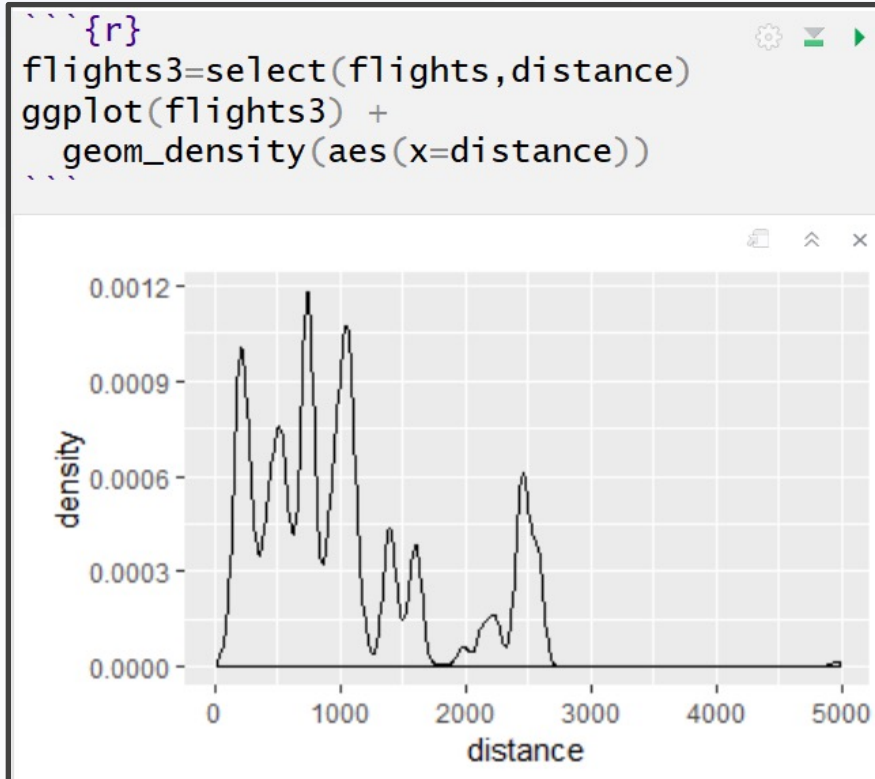
5:17AM

5:33AM

5:42AM

mutate()

- Plethora of Examples
 - Nonlinear Transformation




mutate()

- Plethora of Examples
 - Offsets: lead() and lag()

```
```{r}
flights5=transmute(flights,
 dep_delay,
 lag1_dep_delay=lag(dep_delay))
flights5
```
```

| dep_delay
<dbl> | lag1_dep_delay
<dbl> |
|--------------------|-------------------------|
| 2 | NA |
| 4 | 2 |
| 2 | 4 |
| -1 | 2 |
| -6 | -1 |
| -4 | -6 |



mutate()

- Plethora of Examples
 - Cumulative and Rolling Aggregates

- cumsum()
- cumprod()
- cummin()
- cummax()
- cummean()

```
```{r}
flights6<-transmute(filter(flights,origin=="LGA",
 dest=="CLE",carrier=="UA"),dep_delay,
 rollsum_dep_delay=cumsum(dep_delay))
flights6
```
```

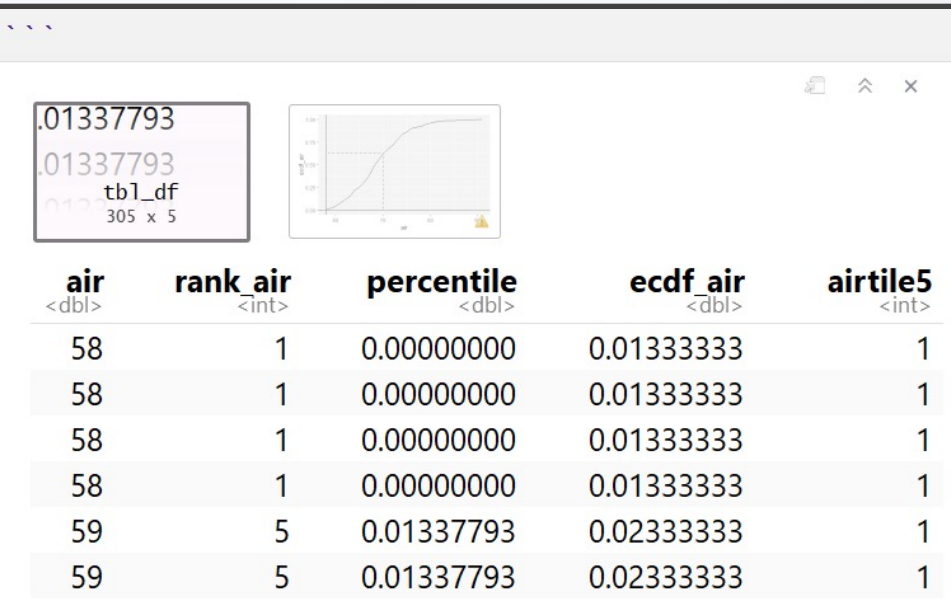
| dep_delay
<dbl> | rollsum_dep_delay
<dbl> |
|--------------------|----------------------------|
| 0 | 0 |
| -1 | -1 |
| 4 | 3 |
| 3 | 6 |
| -6 | 0 |
| -5 | -5 |

mutate()

- Plethora of Examples
 - Ranking
 - min_rank()
 - percent_rank()
 - cume_dist()
 - ntile()

```
flights7<-arrange(transmute(filter(flights,  
                                origin=="LGA",dest=="CLE",  
                                carrier=="UA"),air=air_time,  
                                rank_air=min_rank(air_time),  
                                percentile=percent_rank(air_time),  
                                ecdf_air=cume_dist(air_time),  
                                airtile5=ntile(air,5)),  
air)
```

flights7



The screenshot shows a RStudio window with a data table and a plot. The data table has 6 rows and 5 columns: air, rank_air, percentile, ecdf_air, and airtile5. The plot shows a cumulative distribution function (CDF) for the 'air' variable, with a y-axis labeled 'P(X ≤ x)' and an x-axis labeled 'air'.

| air
<dbl> | rank_air
<int> | percentile
<dbl> | ecdf_air
<dbl> | airtile5
<int> |
|--------------|-------------------|---------------------|-------------------|-------------------|
| 58 | 1 | 0.00000000 | 0.01333333 | 1 |
| 58 | 1 | 0.00000000 | 0.01333333 | 1 |
| 58 | 1 | 0.00000000 | 0.01333333 | 1 |
| 58 | 1 | 0.00000000 | 0.01333333 | 1 |
| 59 | 5 | 0.01337793 | 0.02333333 | 1 |
| 59 | 5 | 0.01337793 | 0.02333333 | 1 |

Information

- Tutorial 3
 - Practice
 - `filter()`
 - `arrange()`
 - `select()`
 - `mutate()`
 - Introduced
 - Piping `%>%`
 - `group_by()`
 - `summarize()`